

# **Session 13:**

## **Supervised learning, part 2**

*Andreas Bjerre-Nielsen*

# Agenda

1. model building
2. model selection
  - cross validation
  - tools for selection

# Vaaaamos

```
In [15]: import warnings
from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings(action='ignore', category=ConvergenceWarning)


import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
```

# Supervised problems (1)

*What is the tradeoff for making supervised regression models?*

```
In [63]: f_bias_var['regression'][2]
```

Out[63]:



## Supervised problems (2)

*What was a remedy to overfitting in linear models? How do we measure overfitting?*

Regularization

- Too many irrelevant features - solved by L1 regularization ~ lasso
- Exploding coefficients - solved by L2 regularization ~ ridge

**Model building**

# Model pipelines (1)


*Is there a smart way to build ML models?*

Yes, we build a pipeline:

- Preprocessing data
  - Standard: adding polynomials, imputation, rescaling
  - Unsupervised learning (more info..)
- Supervised learning

## Model pipelines (2)

*How does the pipeline look?*





## Model pipelines (3)

*What are the advantages of using a pipeline?*

- Ensures good practice - we only fit on training data.
  - No leakage of data from train to test!
- Much less code!

# Applying a model pipeline (1)

*What would this look like in Python?*

```
In [7]: from sklearn.pipeline import make_pipeline
        from sklearn.preprocessing import PolynomialFeatures, StandardScaler

        pipe_preproc = make_pipeline(PolynomialFeatures(),
                                      StandardScaler())

        print(pipe_preproc.steps[0])
        print(pipe_preproc.steps[1])

('polynomialfeatures', PolynomialFeatures(degree=2, include_bias=True, interaction_on
ly=False))
('standardscaler', StandardScaler(copy=True, with_mean=True, with_std=True))
```

## Applying a model pipeline (2)

*Does this remind you of something?*

**YES!**

**Method chaining from Pandas**

## Applying a model pipeline (3)

*Let's some load Boston house price data*

```
In [5]: from sklearn.datasets import load_boston
boston = load_boston()
print(boston['DESCR'])
# print('\n'.join(load_boston()['DESCR'].split('\n')[12:26]))
```

## Applying a model pipeline (4)

*And how do I apply the pipe on the data?*

```
In [9]: pipe_preproc = make_pipeline(PolynomialFeatures(),  
                                     StandardScaler()) # apply preproc - fit on train  
pipe_preproc.fit(X_train) # fit to training  
X_train_prep = pipe_preproc.transform(X_train) # transform training data  
X_test_prep = pipe_preproc.transform(X_test) # transform test data
```

## Applying a model pipeline (5)

*What would it look like if we did use the pipe..?*

The more steps we have, the more code we save.

```
In [10]: poly_trans = PolynomialFeatures()
          scaler = StandardScaler()

# we call both transformations twice on both test and train
X_train_poly = poly_trans.fit_transform(X_train)
X_test_poly = poly_trans.fit_transform(X_test)
scaler.fit(X_train_poly)
X_train_prep_alt = scaler.transform(X_train_poly)
X_test_prep_alt = scaler.transform(X_test_poly)
```

**Model selection**

# Measuring the problem

*Does machine learning work out of the box?*

- In some cases ML works quite well out of the box.
- Often ML requires making careful choices.
  - Note that automated machine learning packages and services exist.
  - E.g. AutoML - this a hot research topic

*Which choices are to be made?*

- We need to pick model building **hyperparameters**.
- E.g.  $\lambda$  for Lasso, Ridge.



# Model validation (1)

*How do we measure our model's performance for different hyperparameters?*


- Remember we cannot use the test set.

*Could we somehow mimick what we do with test data?*

- Yes, we can split the remaining non-test data into training and validation data:
  - we train model for various hyperparameters on training data;
  - pick the hyperparameters which performs best on validation data.

## Model validation (2)

*The non-test data is split into training and validation*



## Model validation (3)

*What would this look like in Python?*

```
In [94]: # splitting into development (2/3) and test data (1/3)
X_dev, X_test, y_dev, y_test = train_test_split(X, y, test_size=1/3, random_state=1)

# splitting development into train (1/3) and validation (1/3)
X_train, X_val, y_train, y_val = train_test_split(X_dev, y_dev, test_size=1/2, random_state=1)
```

## Model validation (4)

Let's train a linear regression model

```
In [95]: from sklearn.linear_model import Lasso, LinearRegression
pipe_lr = make_pipeline(PolynomialFeatures(include_bias=False),
                        StandardScaler(),
                        LinearRegression())

pipe_lr.fit(X_dev, y_dev)
```

```
Out[95]: Pipeline(memory=None,
                  steps=[('polynomialfeatures', PolynomialFeatures(degree=2, include_bias=False, i
nteraction_only=False)), ('standardscaler', StandardScaler(copy=True, with_mean=True,
with_std=True)), ('linearregression', LinearRegression(copy_X=True, fit_intercept=Tru
e, n_jobs=None,
                  normalize=False))])
```

## Model validation (5)

Let's find the Lasso model which performs best in the validation set

```
In [96]: from sklearn.metrics import mean_squared_error as mse

perform = []
lambdas = np.logspace(-4, 4, 33)
for lambda_ in lambdas:
    pipe_lasso = make_pipeline(PolynomialFeatures(include_bias=False),
                               StandardScaler(),
                               Lasso(alpha=lambda_, random_state=1))
    pipe_lasso.fit(X_train, y_train)
    y_pred = pipe_lasso.predict(X_val)
    perform.append(mse(y_pred, y_val))

hyperparam_perform = pd.Series(perform, index=lambdas)

optimal = hyperparam_perform.nsmallest(1)
print('Optimal alpha:', optimal.index[0])
print('Validation MSE: %.3f' % optimal.values[0])
```

Optimal alpha: 0.01  
Validation MSE: 18.421

## Model validation (6)

Let's compare the performance of the Lasso vs. Linear Regression

```
In [97]: pipe_lasso = make_pipeline(PolynomialFeatures(include_bias=False),  
                                   StandardScaler(),  
                                   Lasso(alpha=optimal.index[0]))  
  
pipe_lasso.fit(X_dev,y_dev)  
print('Lasso', round(mse(pipe_lasso.predict(X_test),y_test), 3))  
print('LinReg', round(mse(pipe_lr.predict(X_test),y_test), 3))
```

Lasso 12.399

LinReg 17.802

# Bias and variance (1)

*How do we describe the modelling error?*

From [Wikipedia \(https://en.wikipedia.org/wiki/Bias%E2%80%93variance\\_tradeoff\)](https://en.wikipedia.org/wiki/Bias%E2%80%93variance_tradeoff) yesterday:

- model **bias**: *an error from erroneous assumptions in the learning algorithm*
  - high bias can cause an algorithm to miss the relevant relations between features and target outputs (**underfitting**)
- model **variance**: *an error from sensitivity to small fluctuations in the training set*
  - high variance can cause an algorithm to model the random noise in the training data, rather than the intended outputs (**overfitting**).

## Bias and variance (2)

*So what is overfitting?*

Overfitting is: low bias / high variance

- training our model captures all patterns but we also find some irrelevant
- reacts too much to training sample errors
  - some errors are just noise, and thus we find too many spurious relations
- examples of causes:
  - too much polynomial expansion of variables (`PolynomialFeatures`)
  - non-linear/logistic without properly tuned hyperparameters:
    - Decision Trees, Support Vector Machines or Neural Networks



## Bias and variance (3)

*So what is underfitting?*

Underfitting is: high bias / low variance

- oversimplification of models, cannot approximate all patterns found
- examples of causes:
  - linear and logistic regression (without polynomial expansion)

## Bias and variance (4)

*Not so fast.. OLS is unbiased, right?*

Yes, OLS is unbiased. But...?

- But .. only by assumption..
  - Requires we know the true form of the model.
    - However, we never know do..

*What happens if we introduce regularization?*

- Then model is no longer unbiased.
  - (if we assume the model is true)

## Smarter validation

*Is this approach the smartest way for deciding on choice of hyperparameters?*

**NO**

Our model choice depends a lot on which sample we pick. Could we use more of the data?

# Cross validation

# The holdout method

*How do we get the most out of the data?*

We reuse the data in the development set repeatedly

- We test on all the data
- Rotate which parts of data is used for test and train.

# Leave-one-out CV

*How do we get the most of the data?*

The most robust approach

- Each single observation in the training data we use the remaining data to train.
- Makes number of models equal to the number of observations
- Very computing intensive - does not scale! LOOCV

# K fold method (1)

*How do balance computing time vs. overfitting?*

We split the sample into  $K$  even sized test bins.


- For each test bin  $k$  we use the remaining data for training.

Advantages:

- We use all our data for testing.
- Training is done with  $100 - (100/K)$  pct. of the data, i.e. 90 pct. for  $K=10$ .

# K fold method (2)

In K-fold cross validation we average the errors.





## K fold method (3)

*How would we use K-fold cross validation to select our model?*

We compute MSE for every lambda and every fold (nested for loop)

```
In [98]: from sklearn.model_selection import KFold
kfolds = KFold(n_splits=10)

# outer loop: Lambdas
mseCV = []
for lambda_ in lambdas:
    # inner loop: folds
    mseCV_ = []
    for train_idx, val_idx in kfolds.split(X_dev, y_dev):
        # train model and compute MSE on test fold
        pipe_lassoCV = make_pipeline(PolynomialFeatures(degree=3, include_bias=False),
                                     StandardScaler(),
                                     Lasso(alpha=lambda_, random_state=1))
        X_train, y_train, = X_dev[train_idx], y_dev[train_idx]
        X_val, y_val = X_dev[val_idx], y_dev[val_idx]
        pipe_lassoCV.fit(X_train, y_train)
        mseCV_.append(mse(pipe_lassoCV.predict(X_val), y_val))

    # store result
    mseCV.append(mseCV_)

# convert to DataFrame
lambdaCV = pd.DataFrame(mseCV, index=lambdas)
```

# K fold method (4)

Training the model with optimal hyperparameters and compare MSE

```
In [99]: # choose optimal hyperparameters
         optimal_lambda = lambdaCV.mean(axis=1).nsmallest(1)

         # retrain model using optimal hyperparameters
         pipe_lassoCV = make_pipeline(PolynomialFeatures(include_bias=False),
                                     StandardScaler(),
                                     Lasso(alpha=optimal_lambda.index[0], random_state=1))
         pipe_lassoCV.fit(X_dev,y_dev)

         # compare performance
         models = {'Lasso': pipe_lasso, 'Lasso CV': pipe_lassoCV, 'LinReg': pipe_lr}
         for name, model in models.items():
             score = mse(model.predict(X_test),y_test)
             print(name, round(score, 2))
```

Lasso 12.4

Lasso CV 12.4

LinReg 17.8





## K fold method (5)

*What else could we use cross-validation for?*

- Getting more evaluations of our model performance.
- We can cross validate at two levels:
  - Outer: we make multiple splits of test and train/dev.
  - Inner: within each train/dev. dataset we make cross validation to choose hyperparameters

# Learning curves (1)

*What does a balanced model look like?*





## Learning curves (2)

```
In [52]: from sklearn.model_selection import learning_curve

train_sizes, train_scores, test_scores = \
    learning_curve(estimator=pipe_lasso,
                  X=X_train,
                  y=y_train,
                  train_sizes=np.arange(0.2, 1.05, .05),
                  scoring='neg_mean_squared_error',
                  cv=3)

mse_ = pd.DataFrame({'Train':-train_scores.mean(axis=1),
                    'Test':-test_scores.mean(axis=1)})\
    .set_index(pd.Index(train_sizes,name='sample size'))
print(mse_.head(3))
```


	Train	Test
sample size		
40	0.836417	67.617824
50	1.343929	103.111180
60	1.592246	52.207788

# Learning curves (3)

```
In [53]: f_learn, ax = plt.subplots(figsize=(10,4))
ax.fill_between(train_sizes,
               -test_scores.mean(1) + test_scores.std(1)*1.96,
               -test_scores.mean(1) - test_scores.std(1)*1.96,
               alpha=0.25, label='Train', color='blue')

ax.fill_between(train_sizes,
               -train_scores.mean(1) + train_scores.std(1)*1.96,
               -train_scores.mean(1) - train_scores.std(1)*1.96,
               alpha=0.25, label='Test', color='orange')
ax.set_ylabel('Mean squared error')
ax.set_yscale('log')
ax.legend()
```


Out[53]: <matplotlib.legend.Legend at 0x23991865ef0>



## Learning curves (4)

In [54]: `f_learn`

Out[54]:





# **Tools for model selection**

# Validation curves (1)

```
In [55]: from sklearn.model_selection import validation_curve

train_scores, test_scores = \
    validation_curve(estimator=pipe_lasso,
                    X=X_train,
                    y=y_train,
                    param_name='lasso__alpha',
                    param_range=lambdas,
                    scoring='neg_mean_squared_error',
                    cv=3)

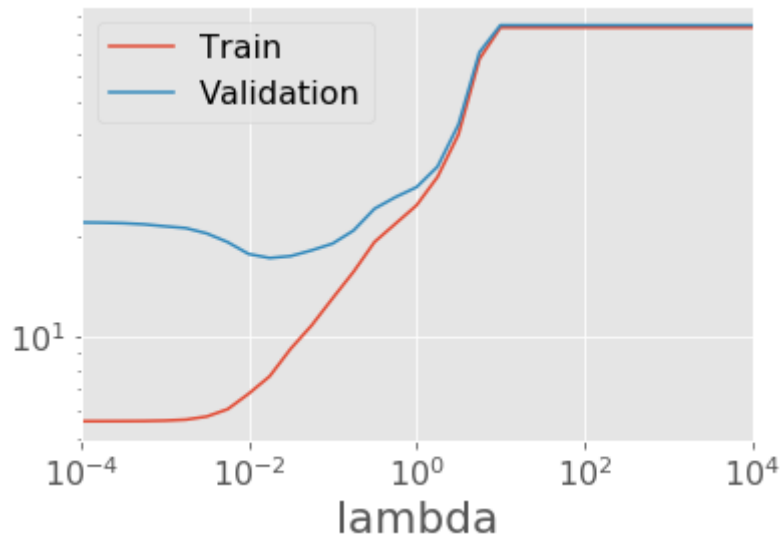
mse_score = pd.DataFrame({'Train':-train_scores.mean(axis=1),
                        'Validation':-test_scores.mean(axis=1),
                        'lambda':lambdas})\
    .set_index('lambda')
print(mse_score.Validation.nsmallest(1))
```

```
lambda
0.017783    17.207944
Name: Validation, dtype: float64
```

## Validation curves (2)

In [56]: `mse_score.plot(logx=True, logy=True)`

Out[56]: `<matplotlib.axes._subplots.AxesSubplot at 0x23991b08208>`



# Grid search (1)

*How do we search for two or more optimal parameters? (e.g. elastic net)*

- Goal: find the optimal parameter combination:

$$\lambda_1^*, \lambda_2^* = \arg \min_{\lambda_1, \lambda_2} MSE^{CV}(X_{train}, y_{train})$$

- Option 1: We can loop over the joint grid of parameters.
  - One level for each parameter.
  - Caveats:
- Option 2: sklearn has GridSearchCV has a tool which tests all parameter combinations.

## Grid search (2)

*How does this look in Python?*

```
In [124]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import ElasticNet
pipe_el = make_pipeline(PolynomialFeatures(include_bias=False),
                        StandardScaler(),
                        ElasticNet())
gs = GridSearchCV(estimator=pipe_el,
                  param_grid={'elasticnet__alpha':np.logspace(-4,4,10)*2,
                              'elasticnet__l1_ratio':np.linspace(0,1,10)},
                  scoring='neg_mean_squared_error',
                  n_jobs=4,
                  iid=False,
                  cv=10)
```

- Notation: double underscore between estimator and hyperparameter, e.g. 'est\_\_hyperparam'
- Scoring: negative MSE as we're maximizing the score ~ minimize MSE.

## Grid search (3)

*What does the grid search yield?*

```
In [ ]: models['ElasticNetCV'] = gs.fit(X_train, y_train)
        for name, model in models.items():
            score = mse(model.predict(X_test), y_test)
            print(name, round(score, 2))
        print()
        print('CV params:', gs.best_params_)
```

## Grid search (4)

*What if we have 10,000 parameter combinations?*

- Option 1: you buy a cluster on Amazon, learn how to parallelize across computers.
- Option 2: you drop some of the parameter values
- Option 3: `RandomizedSearchCV` searches a subset of the combinations.

# The end

[Return to agenda](#)

```
In [2]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import requests
import seaborn as sns

plt.style.use('ggplot')
%matplotlib inline

SMALL_SIZE = 16
MEDIUM_SIZE = 18
BIGGER_SIZE = 20

plt.rc('font', size=SMALL_SIZE)          # controls default text sizes
plt.rc('axes', titlesize=SMALL_SIZE)     # fontsize of the axes title
plt.rc('axes', labelsiz=MEDIUM_SIZE)    # fontsize of the x and y labels
plt.rc('xtick', labelsiz=SMALL_SIZE)     # fontsize of the tick labels
plt.rc('ytick', labelsiz=SMALL_SIZE)     # fontsize of the tick labels
plt.rc('legend', fontsize=SMALL_SIZE)    # legend fontsize
plt.rc('figure', titlesize=BIGGER_SIZE)  # fontsize of the figure title

plt.rcParams['figure.figsize'] = 10, 4 # set default size of plots
```